

CMSC 27200 - Guerrilla Section 6.

TOPICS:

1. Network flow.

- Ford-Fulkerson
- Augmenting paths and residual graphs.
- Remarks on running time.

2. Flow reductions.

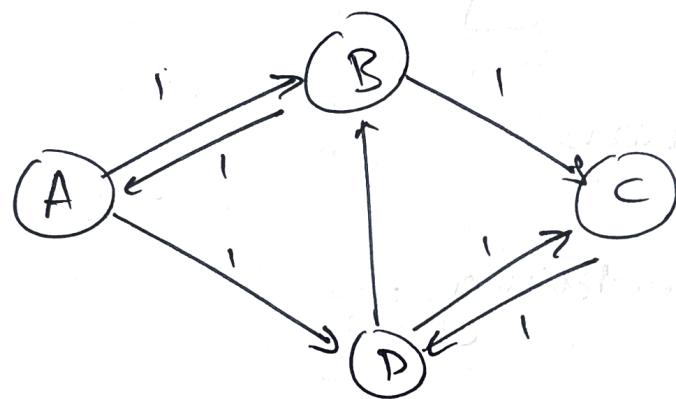
- Bipartite matching:

Linear Programming & Network Flow

So what's the Spark notes version of LPs and network flow?

⇒ Network flow.

Suppose you have a computer network.



Each edge is labeled w/ a maximum amt of data that can be sent b/w pairs of computers (bandwidth).

Now you pick a source vector A and sink vector B.

Your goal is to send the maximum amount of data b/w source and sink.

↳ This is called maximum flow.

How do you model this?

Defn: Given $G = (V, E)$ directed and st ~~UV~~.

capacities $c_e > 0 \quad \forall e \in E$ and source-sink pair $S, t \in V$, an $s-t$ -flow in G is

a collection of values $\{f_e\}_{e \in E}$ ~~st~~.

associated w/ each edge e .

1) It satisfies capacity constraints: $\forall e \in E$

$$0 \leq f_e \leq c_e.$$

2) Flow is conserved at each vertex except the source and sink.

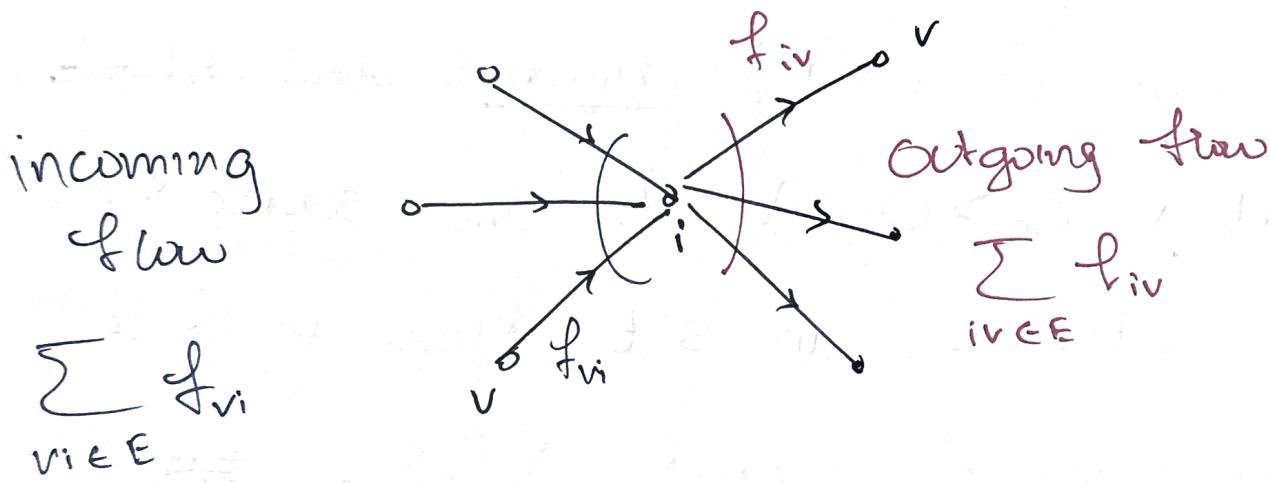
(140)

$\forall i \in V \setminus \{s, t\}$ we have

$$\sum_{(v,i) \in E} f_{vi} = \sum_{(i,v) \in E} f_{iv}.$$

flow going into i = flow leaving i .

Pictorially...



Flow conservation \equiv "in-flow = out-flow"

And now we can state the problem precisely

the maximum flow problem (aka "single-commodity maximum flow") is defined as...

(111)

INPUT: A directed graph $G = (V, E)$. edge

capacities $c_e > 0 \quad \forall e \in E$, source vertex

$s \in V$ and sink vertex $t \in V$

Output: An s - t flow of G^* w/ maximum

size ~~is~~. where

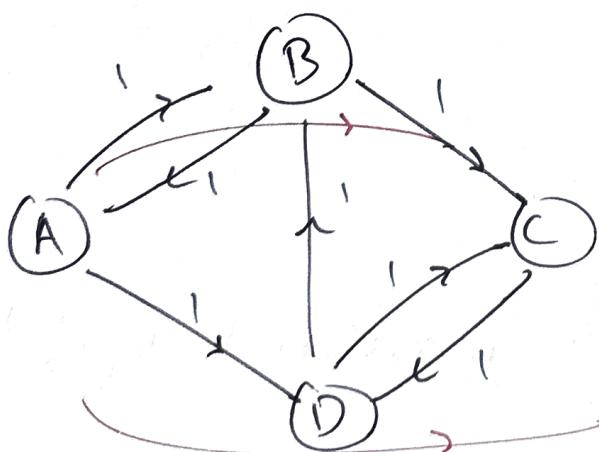
$$\text{max size}(f) := \max \sum$$

$$\text{size}(f) := \sum_{(s, v) \in E} f_{sv}$$

amt of flow leaving s .

Example: What's a maxflow of this graph?

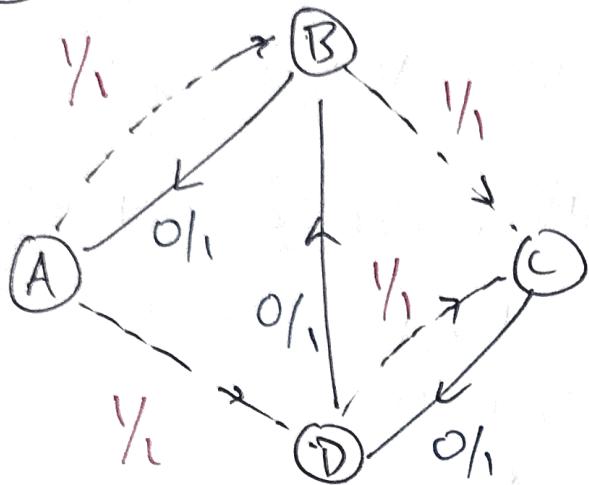
w/ $s = A$, $t = C$.



send one unit b/w $A \rightarrow B \rightarrow C$

another b/w $A \rightarrow D \rightarrow C$.

142



Edge	Flow
AB	1
BA	0
BC	1
AD	1
DB	0
DC	-1
CD	0

$$\text{Size}(f) = \sum_{(s,i) \in E} f_{si} = f_{AB} + f_{AD} = 1 + 1 = 2.$$

→ Questions?

Sparknode punchlines:

⇒ Sparknode punchlines.

= ~~How do you solve~~

These are two questions to ask:

1. How do you compute a maxflow (algo?)

2. How do you know you've computed a maxflow?

3. What can I use maxflow to solve?

→ Item #1: how do I compute maxflow?

- Ford - Fulkerson: on first pass this can be described as the following method.

1) Start w/ zero flow.

2) Repeat: find an s-t path w/ nonzero capacity, and push as much flow as possible.

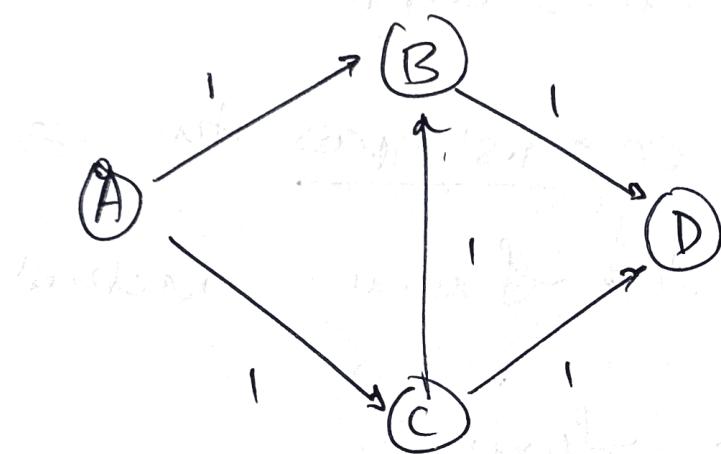
But this is not an algorithm!

→ we didn't specify how to find the paths.

(44)

→ If you're not careful, this method won't find a maxflow

↪ Good example: graphs that look like.

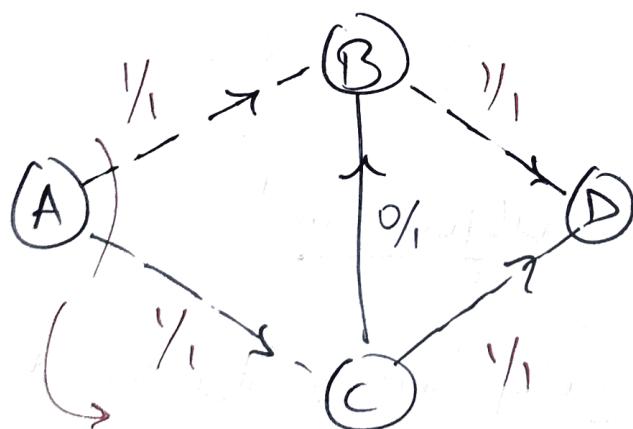


$$s = A$$

$$t = D$$

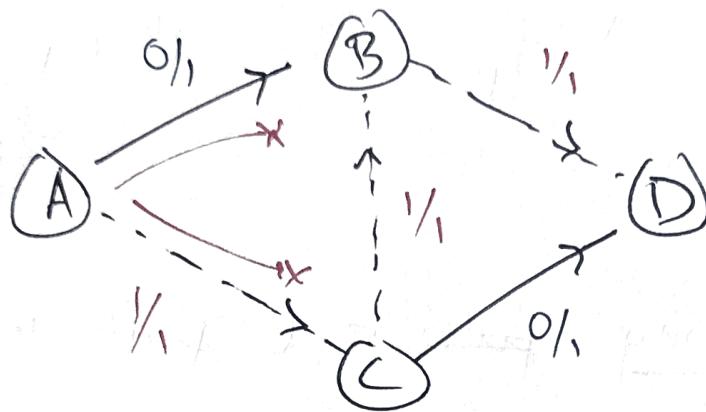
If Ford-Fulkerson picks paths A-B-D

then A-C-D then you get a maxflow.



Size = 2.

But if Ford-Fulkerson picks A - C - B - D then
you're toast.



There's no A \rightarrow D path w/ available capacity

And you've only computed a size = 1 flow

→ The fix is to use the residual graph
and find augmenting paths

Defn: Given a directed graph $G = (V, E)$ and

w/ capacities $C_e > 0$ and s-t flow f , the

residual graph is $G_f = (V, E_f)$. ~~and~~

~~note~~ w/ residual capacities. $C_e^f > 0$ $\forall e \in E_f$.

46

$$C_{ij}^f = \begin{cases} C_{ij} - f_{ij} & \text{if } (i,j) \in E \text{ and } \\ & f_{ij} \leq C_{ij} \\ f_{ji} & \text{if } (j,i) \in E \text{ and } \\ & f_{ji} > 0. \end{cases}$$

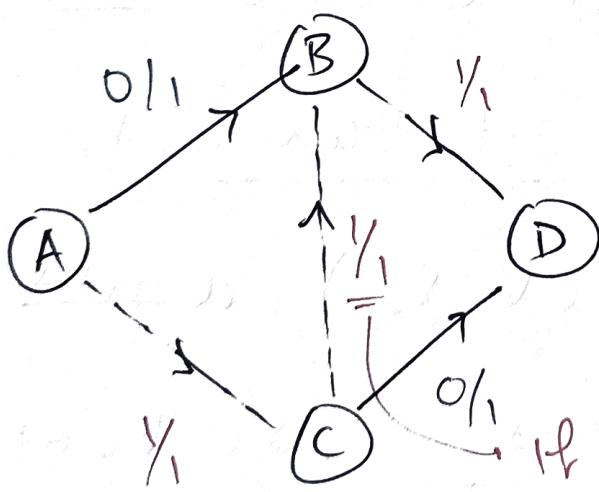
→ An augmenting path $s-t$ path in G^f is

an $s-t$ path in G^f w/ ~~capacity~~ > 0

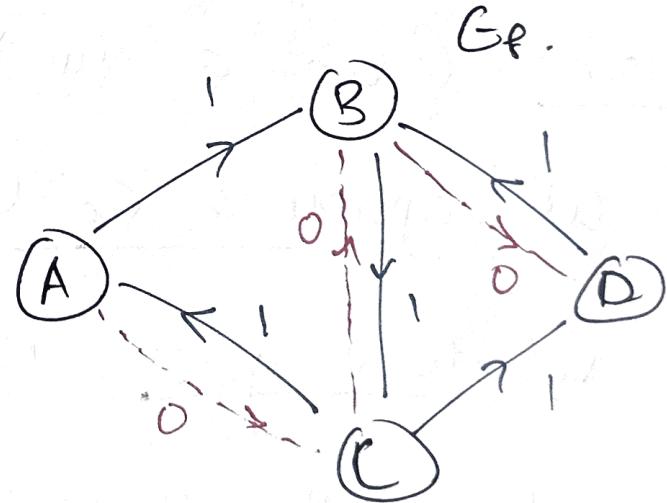
Capacity on each edge

→ What's the intuition?

→ The residual graph tells you how to
"undo" a flow that you pushed.



If we could
"undo" his edge flow

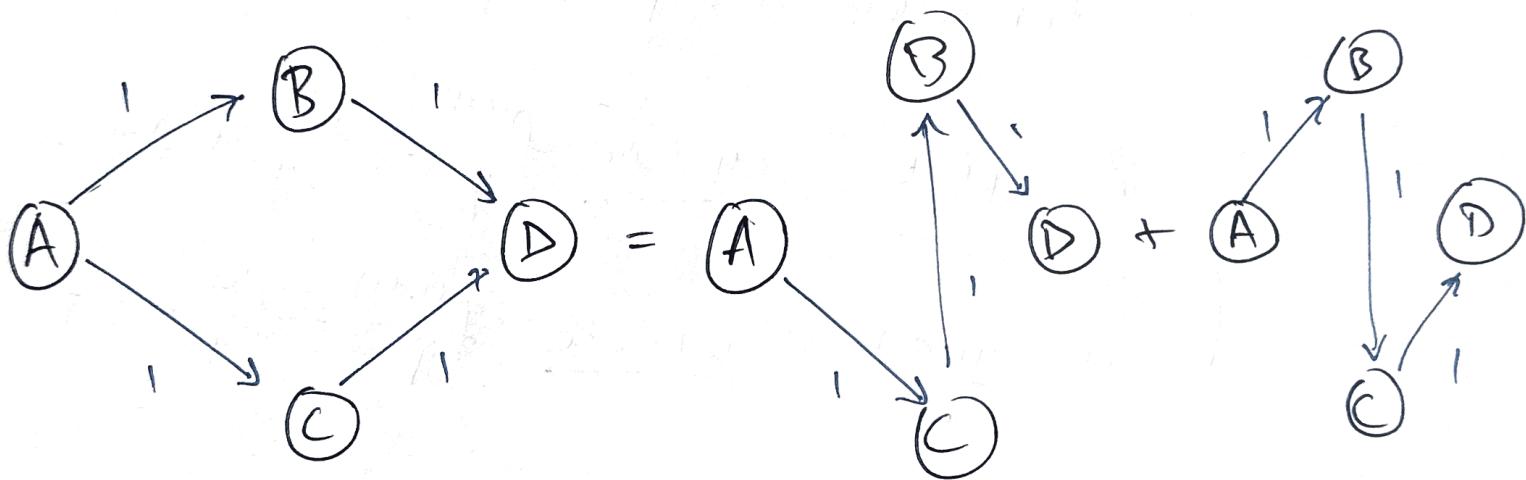


$A \rightarrow B \rightarrow C \rightarrow D$
is an aug path!

along this edge we can get unstuck.

If you route an augment flow along the augmenting path $A \rightarrow B \rightarrow C \rightarrow D$ in G^F then residual graph G^F . Then you under flow sent in the $C \rightarrow B$ direction.

→ Summing the flows ~~across all~~ gives you the max flow.



→ This is Ford-Fulkerson.

Proc: FORD-FULKERSON.

INPUT: $G = (V, E)$ directed, $c_e > 0 \forall e \in E$.

Source $s \in V$, Sink $t \in V$.

(148)

1. Initialize $f_e^{(0)} = 0$. $\forall e \in E$.

2. Do for $t=0, 1, \dots$: ~~until no augmenting~~

- Compute $G^{f^{(t)}}$

- Compute P an augmenting path

- s-t path in $G^{f^{(t)}}$, let flow pushed = Δf .

- If no path: return $f^{(t)}$

- For each $(i, j) \in P$:

$$\left\{ \begin{array}{l} \text{if } ij \in E : \text{update } f_{ij}^{(t+1)} = f_{ij}^{(t)} + \Delta f \\ \text{if } ji \in E : \text{update } f_{ji}^{(t+1)} = f_{ji}^{(t)} - \Delta f \end{array} \right.$$

Do you need to recompute G^{f_t} from scratch

in each iteration? Nope!

→ Maintain a counter on residual capacities.

→ Questions?

→ How do you know 'terminating' = maxflow?

① Duality: mincut = maxflow.

- ~~think~~ If I'm sending data from $S \rightarrow t$.

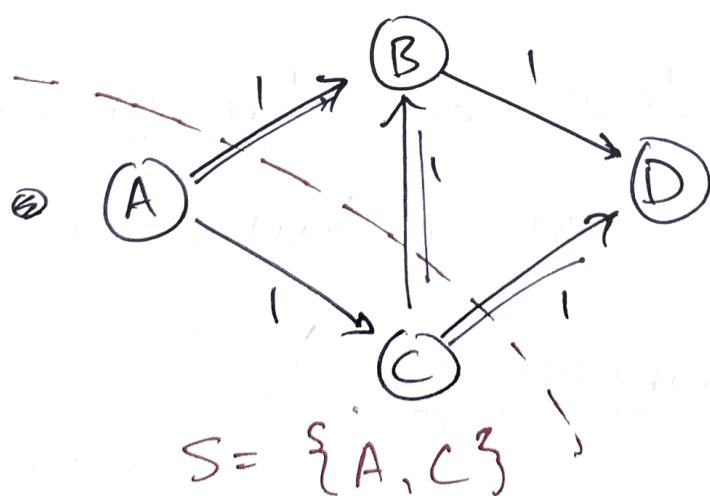
then all the data I send has to cross from one side to another.

↪ Formalize: Let $S \subseteq V$ be a subset

of vertices. (this is called a cut).

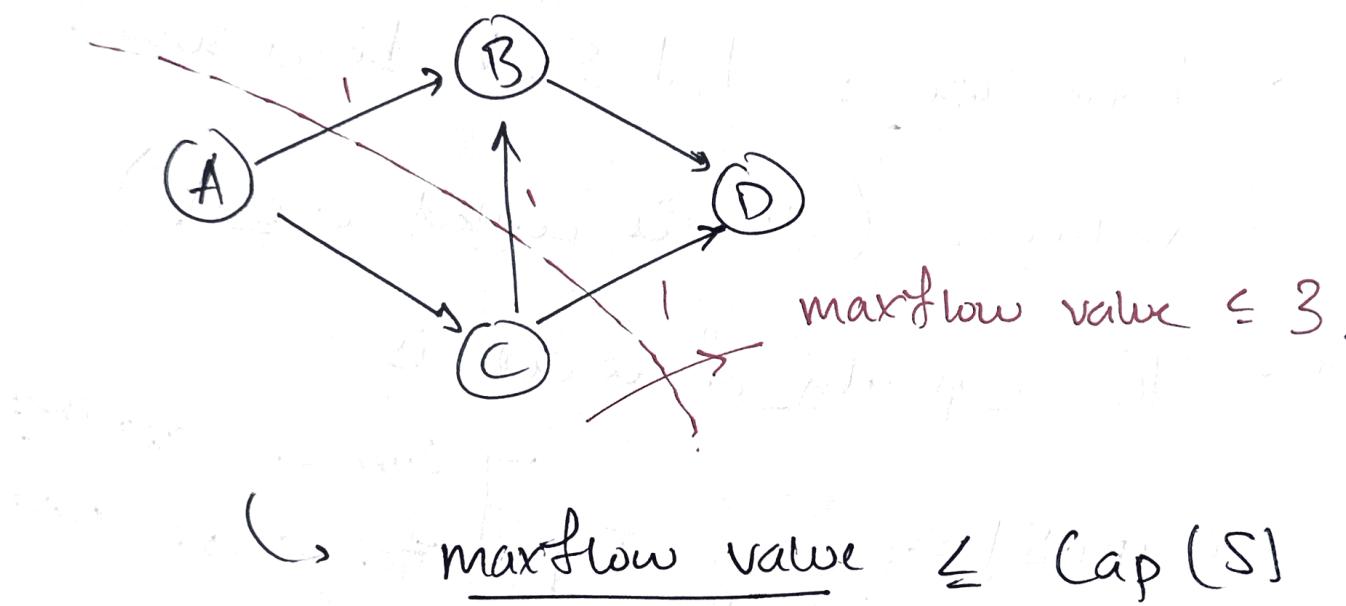
↪ The capacity of a cut is

$$\text{Cap}(S) := \sum_{\substack{i,j \in E : i \in S \\ j \notin S}} c_{ij} \quad \begin{array}{l} \text{sum of } \underline{\text{capacities}} \\ \text{for edges leaving } S. \end{array}$$



$$\text{Cap}(S) = 1 + 1 + 1 = 3.$$

- If the flow is end respects capacity constraints, then the size of the flow can only be at most the capacity of any s-t cut. (i.e. cut S s.t. $i \in S$ and $t \notin S$)



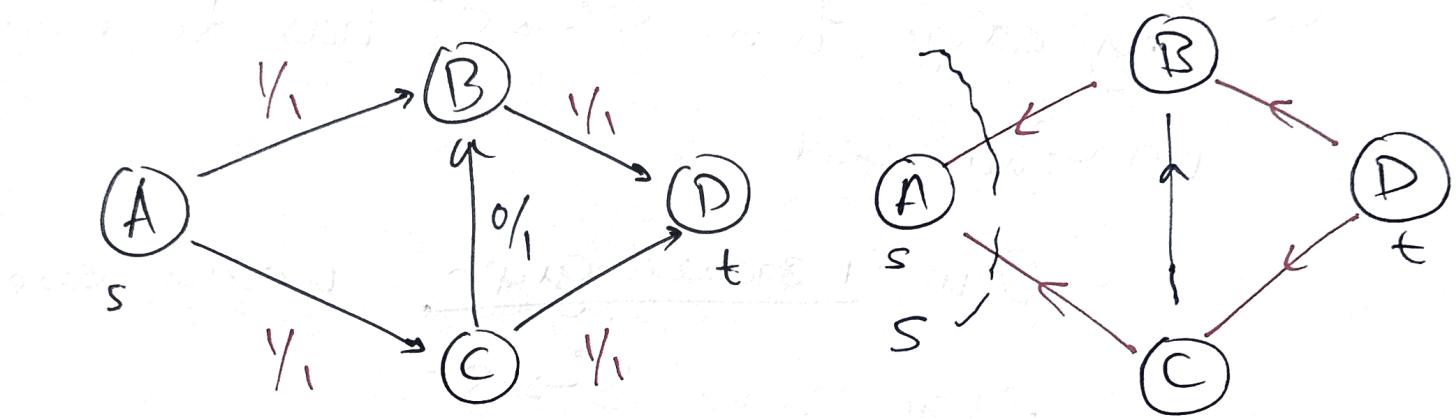
A s-t cuts S EV.

- But now... think of the smallest capacity s-t cut (this is called the minimum cut).

- Let's compare it to the value of the size.

maxflow computed by Ford-Fulkerson

- If Ford-Fulkerson terminates then there are no more augmenting s-t paths in the residual graph.



- Take the set of vertices reachable from s in G^f . This is a cut S^* .

- The capacity of S^* = value of flow computed by Ford-Fulkerson.

\hookrightarrow If not! Then \exists $e \in E$ st.

e leaves S and

↪ Any edge from $S^* \rightarrow \bar{S}^*$ has to be at full capacity

- As the residual graph only has edges going into S^* .

↪ Any edge from $\bar{S}^* \rightarrow S^*$ has to have no flow sent

- O.w. residual graph would have edge from $S^* \rightarrow \bar{S}^*$

- And now what we get is:

$$\text{Cap}(S^*) = \text{Size}(f)$$

- But remember ~~at~~ A s-t cut. S.

~~maxflow~~ $\text{Size}(\text{maxflow}) \leq \text{Cap}(S)$ so

let ~~S~~ be $S = \text{mincut}$.

$$\begin{aligned}
 \text{Cap}(\min \text{-s-t cut}) &\leq \text{Cap}(S^*) \\
 &= \text{Size}(f) \\
 &\leq \text{Size(maxflow)} \\
 &\leq \text{Cap}(\min \text{-s-t cut})
 \end{aligned}$$

→ mincut = maxflow and Ford-Fulkerson

Outputs a maxflow.

→ Questions?

Question: Suppose $e \in E$ is an edge such

that decreasing its

~~Suppose~~ Suppose I pick an edge e

Suppose S is a minimum capacity s-t cut and e is an edge leaving S .

↳ How does the maxflow value change if I decrease $\underline{c_e}$?

(154)

↪ How does the maxflow value change if I increase C_e ?

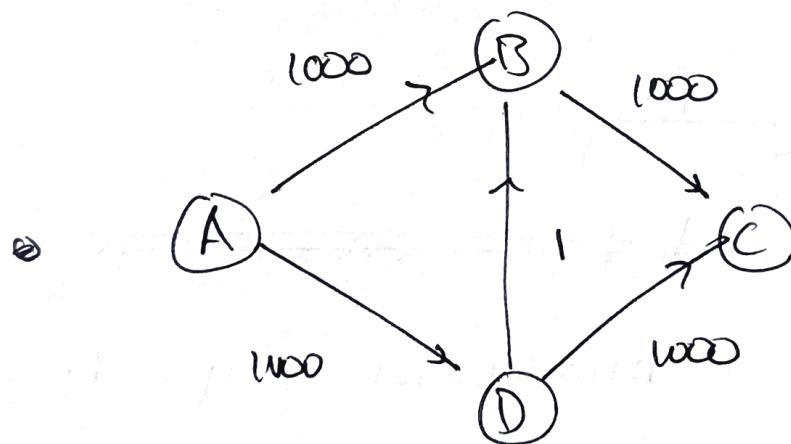
→ Think about it.

(in slides)

⇒ A remark about running time.

Does Ford-Fulkerson always run in polynomial time?

↪ No! Because we're still not careful in specifying what augmenting path to take



give me a sequence of augmenting paths s.t.

it takes 1000 augmenting paths to compute the maxflow.

→ If $U = \max_{e \in E} c_e$. FORD-FULKERSON

might take $O(|E| \cdot U)$ time.

→ Doomed? nah.

↪ If you pick the path of shortest length from s-t then it terminates in time.

$$O(|V| \cdot |E|^2)$$

→ Very active line of research ...

↪ Current fastest approximate maxflow alg

runs in time $\tilde{O}(n \cdot \log^{\epsilon} n \cdot \log \log^6 n \cdot \log(1/\epsilon))$

If $n = |V|$.

↪ Faster than sorting.

↪ Current fastest exact $O(m^{1+\epsilon}) \cdot \lg^c(n \cdot \log m)$

wl $c = O(1)$, $m = |E|$.

(156)

Linear Programming

So what was the bigger

⇒ Using flow to solve other problems.

A key part about learning maxflow is knowing how to reduce other problems to maxflow.

Problem: Bipartite matching

INPUT: A bipartite graph $G = (A \cup B, E)$.

OUTPUT:

⇒ Maxflow reductions.

A key part about learning network flow is also seeing how to use it to solve other problems

Example : Scheduling.

Suppose We're trying to figure out how to assign n TAs to n office hour slots. : ~~and~~

1. Every TA has a certain # of OH slots

that they ~~are~~ are available for

2. Every OH slot needs to be filled by exactly 1 TA.

3. Every TA must be assigned to 1 OH slot.

Given a set of OH slots and TA preferences
can you tell if its possible to schedule the
TAs s.t. the above constraints hold?

→ Let's formalize this.

(158)

INPUT: n TAs $A = \{1, \dots, n\}$ and n OH slots.

$$\{S_1, \dots, S_n\}$$

INPUT: n TAs $A = \{1, \dots, n\}$, n OH slots.

$B = \{S_1, \dots, S_n\}$. and a preference map.

$\sigma : A \rightarrow 2^B$. where if $i \in A$ is a TA then

$\sigma(i) = \{\text{set of OH slots } i \text{ is available for}\}$.

OUTPUT: Yes if \exists an assignment of TAs to

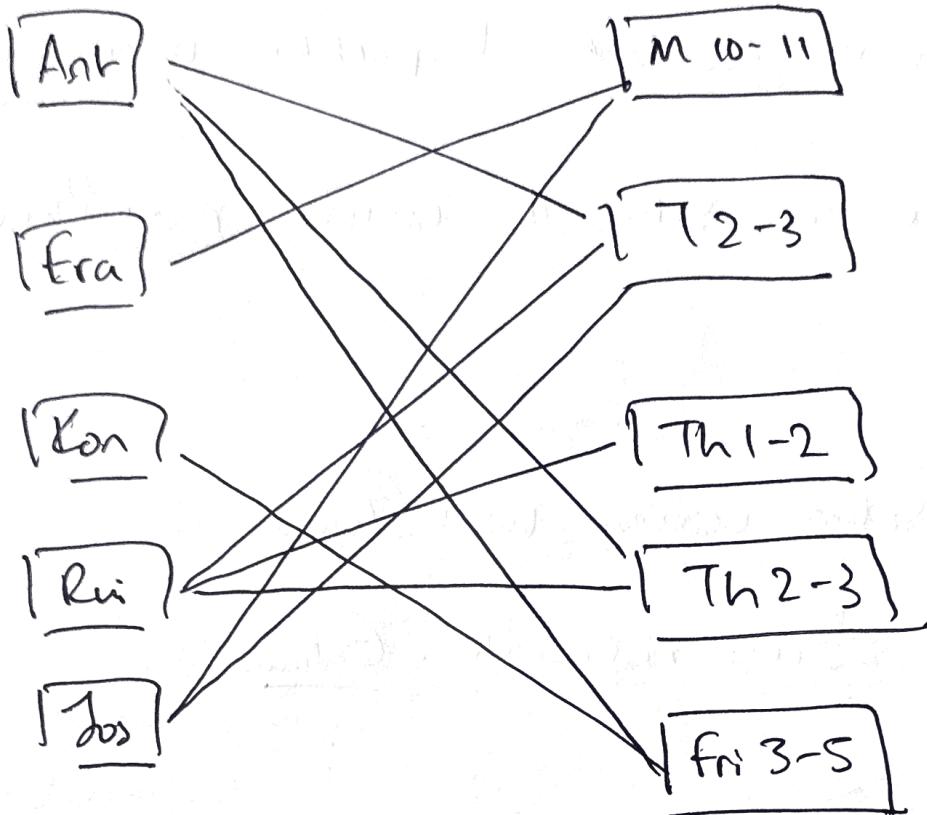
OH sl. st. each TA is paired w/ a unique

OH and vice versa.

for example -- $A = \{\text{Ant, Era, Kan, Rui, Joo}\}$

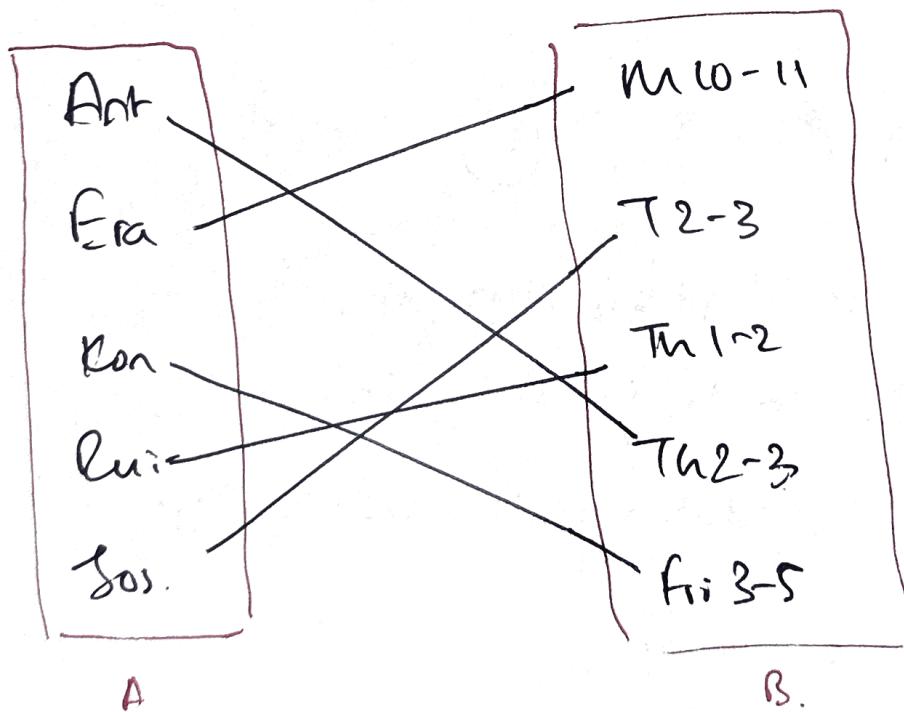
Ant $B = \{\text{Mon 10-11, Tues 2-3, Thurs 1-2, Thurs 2-3, Fri 3-5}\}$

And now the preference map σ .



This is called
a bipartite
graph.

And as it would turn out this is what we come
up w/



This is a
perfect
matching.

This graph we drew (*) is called a bipartite

(60)

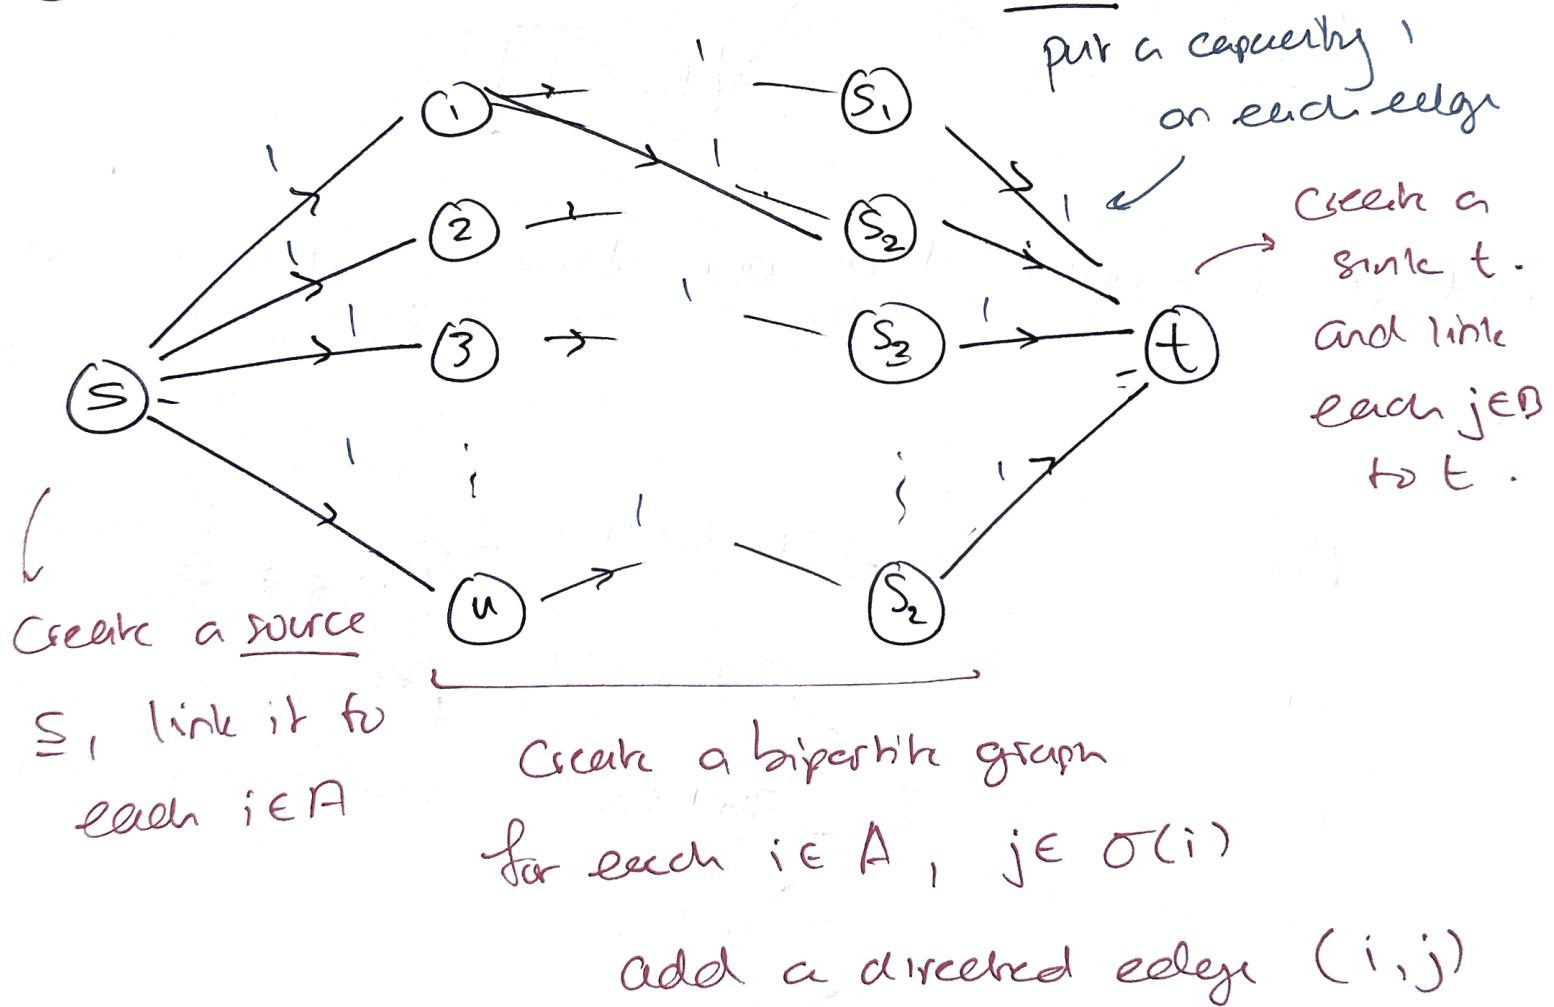
This problem is called bipartite matching.

→ And we can solve it using maxflow.

→ Questions?

⇒ How to solve using maxflow.

① Construct a flow network G_{flow}



- ② State which nodes to route flow b/w
 ↳ Rate a maxflow b/w s and t.
- ③ State what to output for your problem
depending on what maxflow returns
- "If the maxflow ~~water~~ size = n then
return Yes. o.w. return No.

And that's our algorithm.

→ Questions?

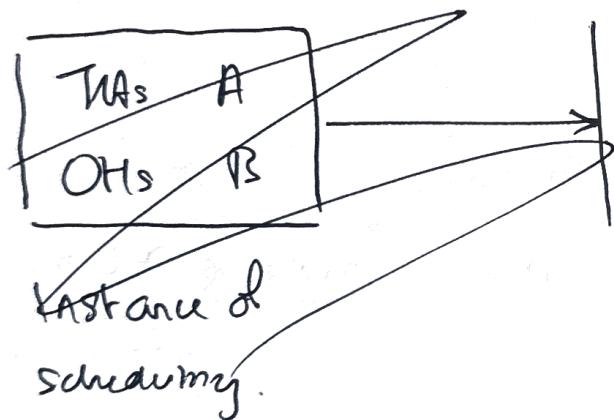
So... what did it mean to "solve this problem using maxflow"?

→ We took an instance of our original problem (scheduling) i.e. all the inputs to scheduling.

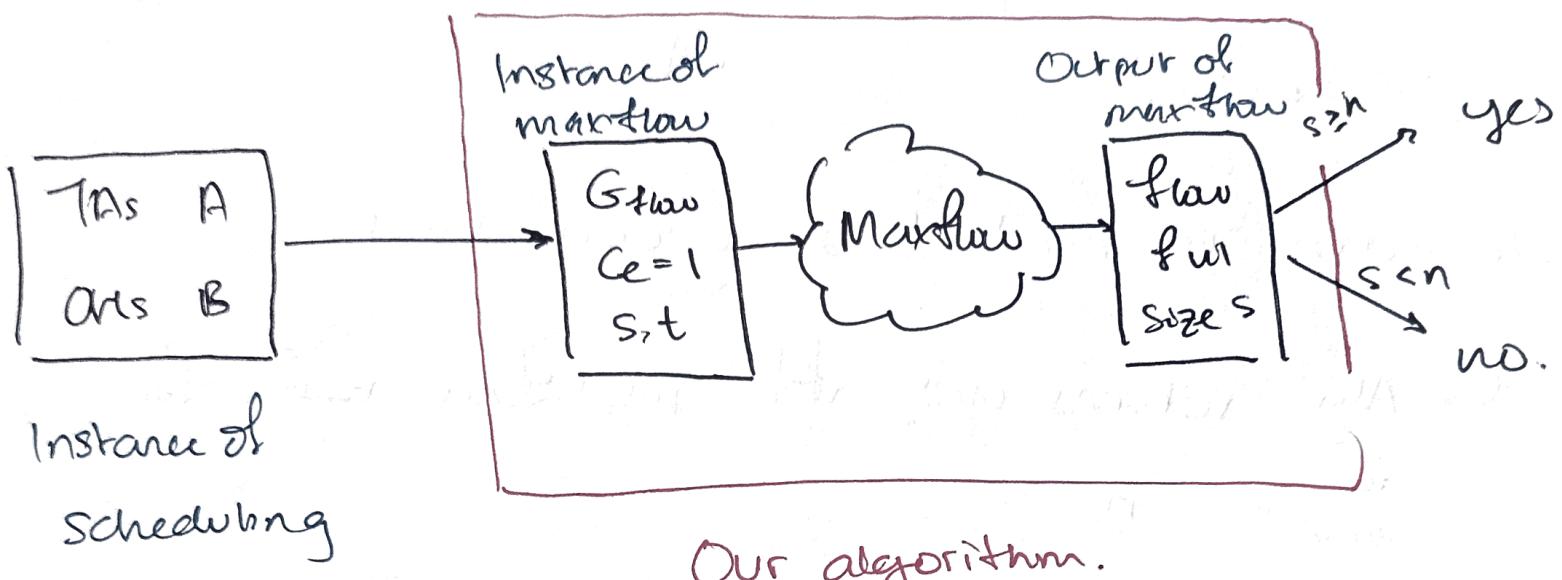
(162)

- We created an instance for maxflow.
i.e. a set of inputs for maxflow.
- We ran maxflow on our new instance.
- Based on the output, we returned a solution for scheduling.

Pictorially.



Pictorially:



→ This is called a "reduction to maxflow"

↳ why? If we have a program to solve maxflow, → we immediately have a program to solve Scheduling.

↳ "Maxflow is at least as hard as Scheduling"

↳ This is how we reason about hardness of problems when we don't have an algo to solve them.

→ In this case we can definitely solve maxflow

(164)

maxflow in polytime so we're good.

→ Questions?

Of course we need to show correctness.

↪ Alg returns yes iff maxflow ^{size} ~~value~~ for G_{flow} is $= n$.

→ We need to show, maxflow ^{size} ~~value~~ for $G_{\text{flow}} = n$ iff. \exists a perfect matching / scheduling

↪ Key point: if all capacities are integer then maxflow value is integer.

~~HTTS:~~: f^* is maxflow and $\text{size}(f^*) = n$.

Pf: "⇒" Suppose size of maxflow = n.

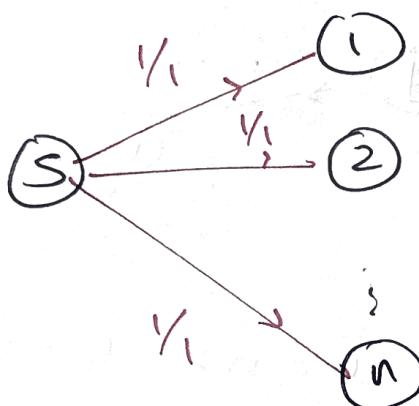
WTS: every TA is assigned and assignment is unique.

→ Every TA is assigned to 1 OH.

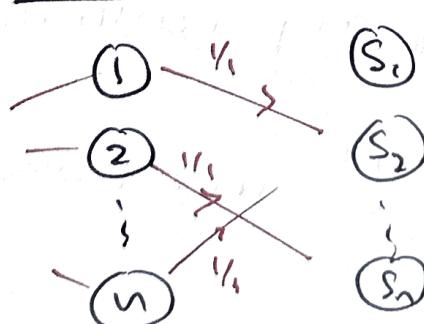
- Note that $\text{size}(f^*) = n$.

↳ flow leaving $S = n$.

↳ By flow conservation, every $i \in A$
has one unit of flow incoming

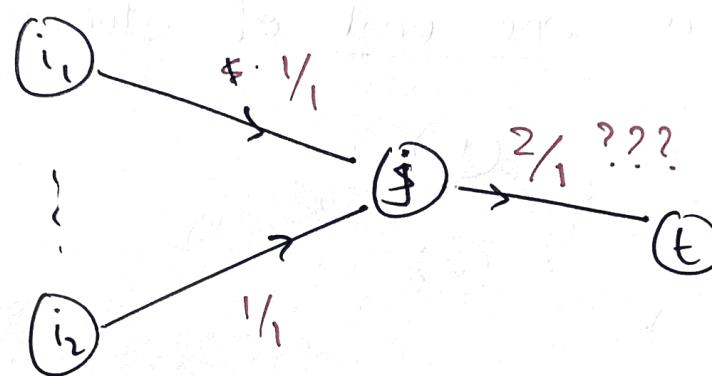


- Matching is integral: → for each $i \in A$, it must send all of its 1 unit of incoming flow along a single edge.
- Because all edges go from A \rightarrow B, all TAs are matched.



→ The assignment is unique i.e. no 2 TAs can match to the same Obj.

- Suppose not. Then $i_1, i_2 \in A$ matched to $\hat{s} \in B$.



- But by defn that means $i_1 \rightarrow s$ sends 1 unit of flow to s , i_2 sends ~~one~~ 1 unit of flow to s .

$$\hookrightarrow \text{Incoming flow to } s = 2.$$

- Outgoing flow = 2 by flow conservation
- Violating capacity constraint on (\hat{s}, t)
- Contradicting f^* was maxflow. \square .

→ Questions?

Other problems that you can solve w/ maxflow?

- Global min cut.
- Vertex capacitated maxflow.

Bigger picture

- Maxflow is a linear program.
- LPs are super ubiquitous. (P-Complete).

Some things you should know how to do w/ an LP.

- Write an LP which models a combinatorial problem.
- Take an LP dual.
- Interpret the dual.
- Compute value of a zero-sum game.